

# RDD FUNDAMENTALS



# INTERACTIVE SHELL

```
ubuntu@ip-10-0-53-24: ~  
ubuntu@ip-10-0-53-24:~$ dse spark  
Welcome to  
  
      /\_/\_/\_/\_/\_/\_/\_/\_/  
     /  V  V  V  V  V  V  V  \  
    /_____\_/_/_/_/_/_/_\_\  
   version 0.9.1  
  
Using Scala version 2.10.3 (Java HotSpot(TM) 64-Bit Server VM, Java 1.7.0_51)  
Type in expressions to have them evaluated.  
Type :help for more information.  
Creating SparkContext...  
Created spark context..  
Spark context available as sc.  
Type in expressions to have them evaluated.  
Type :help for more information.  
  
scala> val myRDD = sc.cassandraTable("tinykeyspace", "keyvaluetable")  
myRDD: com.datastax.bdp.spark.CassandraRDD[com.datastax.bdp.spark.CassandraRow] = Cassan  
draRDD[0] at RDD at CassandraRDD.scala:32  
  
scala> myRDD.count()  
res2: Long = 5  
  
scala>
```


(Scala, Python and R only)

## Driver Program


```
scala> val keyValueRDD = sc.cassandraTable("tinykeyspace", "keyvaluetable")
keyValueRDD: com.datastax.spark.connector.rdd.CassandraRDD[com.datastax.spark.connector.CassandraRow] = CassandraRDD[0] at RDD at CassandraRDD.scala:49

scala> keyValueRDD.count()
res2: Long = 4

scala>
```



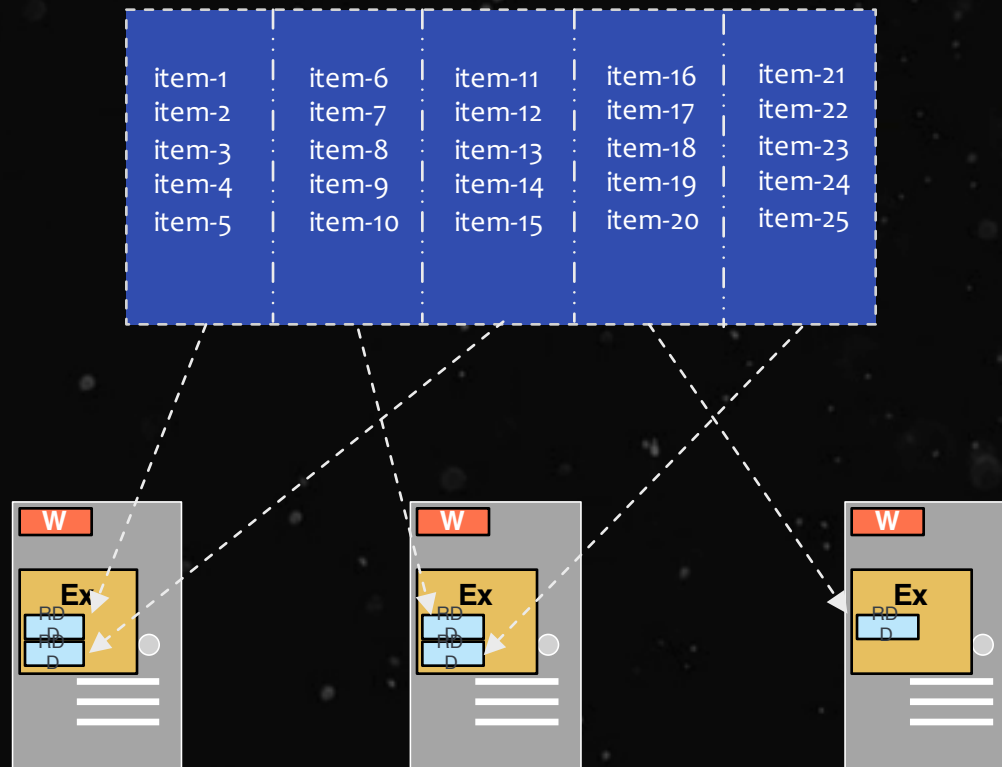
Worker Machine




Worker Machine

*more partitions = more parallelism*

## RDD



## RDD w/ 4 partitions



A base RDD can be created 2 ways:

- Parallelize a collection
- Read data from an external source (S3, C\*, HDFS, etc)

# PARALLELIZE



```
# Parallelize in Python
wordsRDD = sc.parallelize(["fish", "cats", "dogs"])
```

---



```
// Parallelize in Scala
val wordsRDD= sc.parallelize(List("fish", "cats", "dogs"))
```

---



```
// Parallelize in Java
JavaRDD<String> wordsRDD = sc.parallelize(Arrays.asList("fish", "cats", "dogs"));
```

- Take an existing in-memory collection and pass it to SparkContext's parallelize method
- Not generally used outside of prototyping and testing since it requires entire dataset in memory on one machine

## READ FROM TEXT FILE



# Read a local txt file in Python

```
linesRDD = sc.textFile("/path/to/README.md")
```

---

- There are other methods to read data from HDFS, C\*, S3, HBase, etc.



// Read a local txt file in Scala

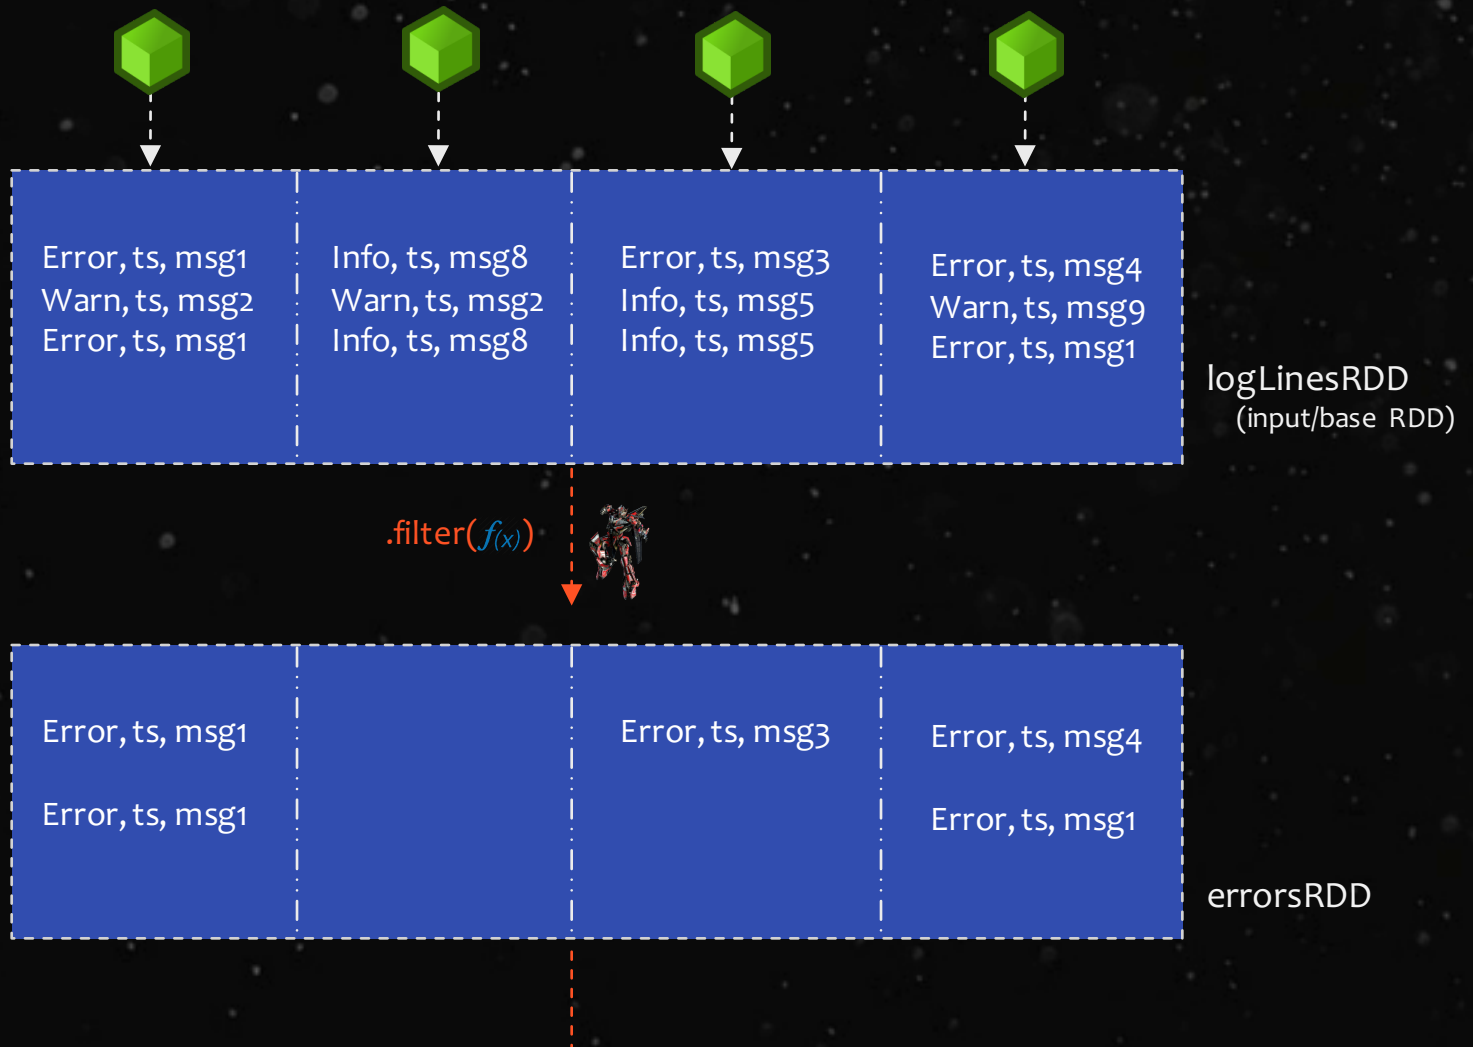
```
val linesRDD = sc.textFile("/path/to/README.md")
```

---




// Read a local txt file in Java

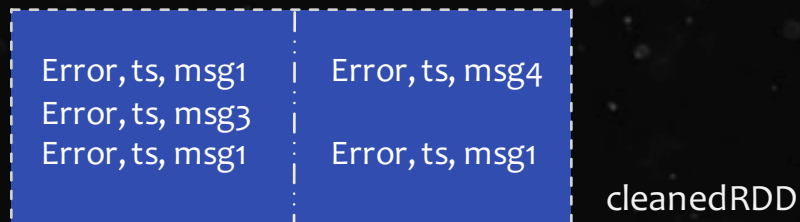
```
JavaRDD<String> lines = sc.textFile("/path/to/README.md");
```







`.coalesce( 2 )`



`.collect( )`



```
scala> val keyValueRDD = sc.cassandraTable("tinykeyspace", "keyvaluestable")
keyValueRDD: com.datastax.spark.connector.rdd.CassandraRDD[com.datastax.spark.connector.CassandraRow] = CassandraRDD[0] at RDD at CassandraRDD.scala:49

scala> keyValueRDD.count()
res2: Long = 4

scala>
```

Driver

Execute DAG!

.collect( )



```
ec2-user@ip-10-0-12-60:~$ spark-shell
[ec2-user@ip-10-0-12-60 ~]$ use spark
Welcome to

Spark version 1.1.0


Using Scala version 2.10.4 (Java HotSpot(TM) 64-Bit Server VM, Java 1.7.0_71)
Type in expressions to have them evaluated.
Type :help for more information.
Creating SparkContext...
Created spark context.
Spark context available as sc.
Type in expressions to have them evaluated.
Type :help for more information.

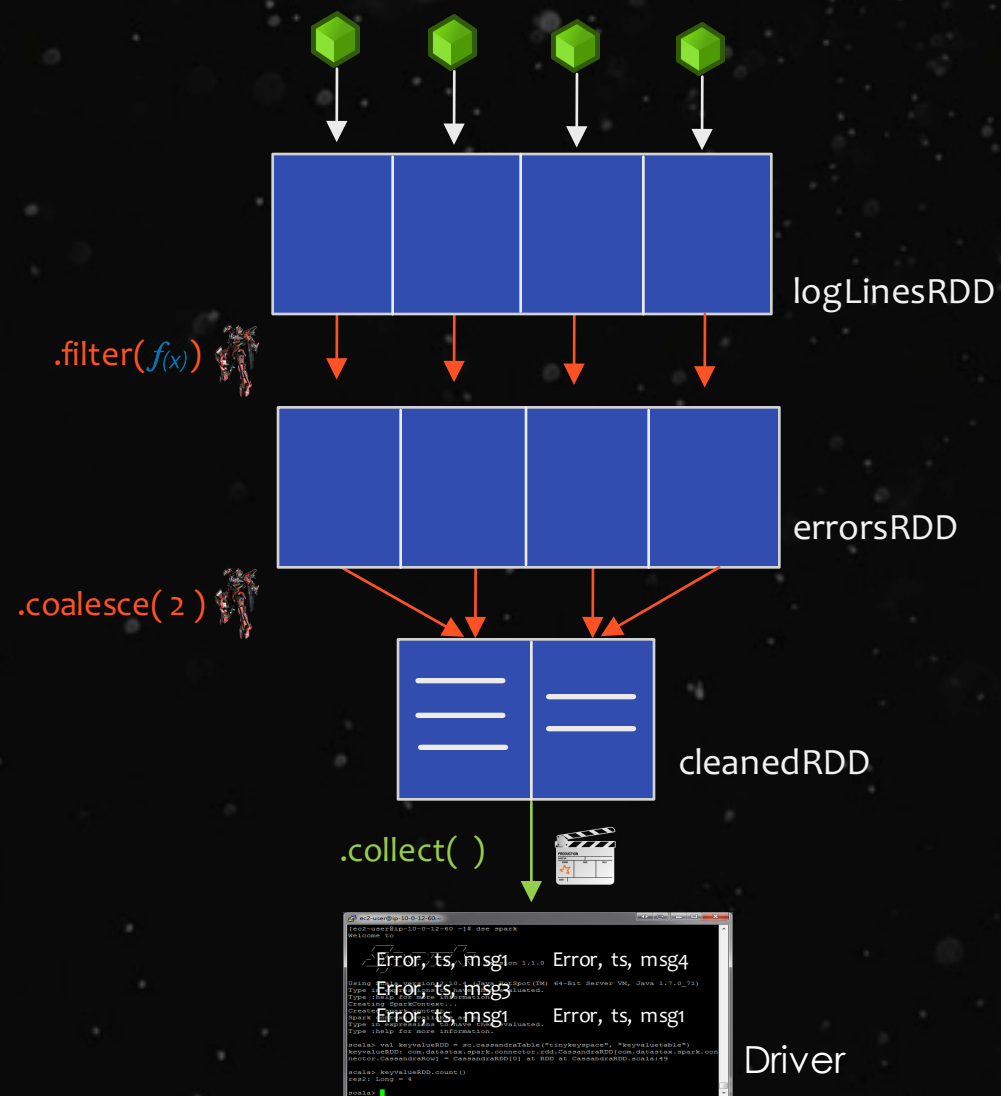
scala> val keyValueRDD = sc.cassandraTable("tinykeyspace", "keyvaluestable")
keyValueRDD: com.datastax.spark.connector.rdd.CassandraRDD[com.datastax.spark.connector.CassandraRow] = CassandraRDD[0] at RDD at CassandraRDD.scala:49

scala> keyValueRDD.count()
res0: Long = 4


scala>
```

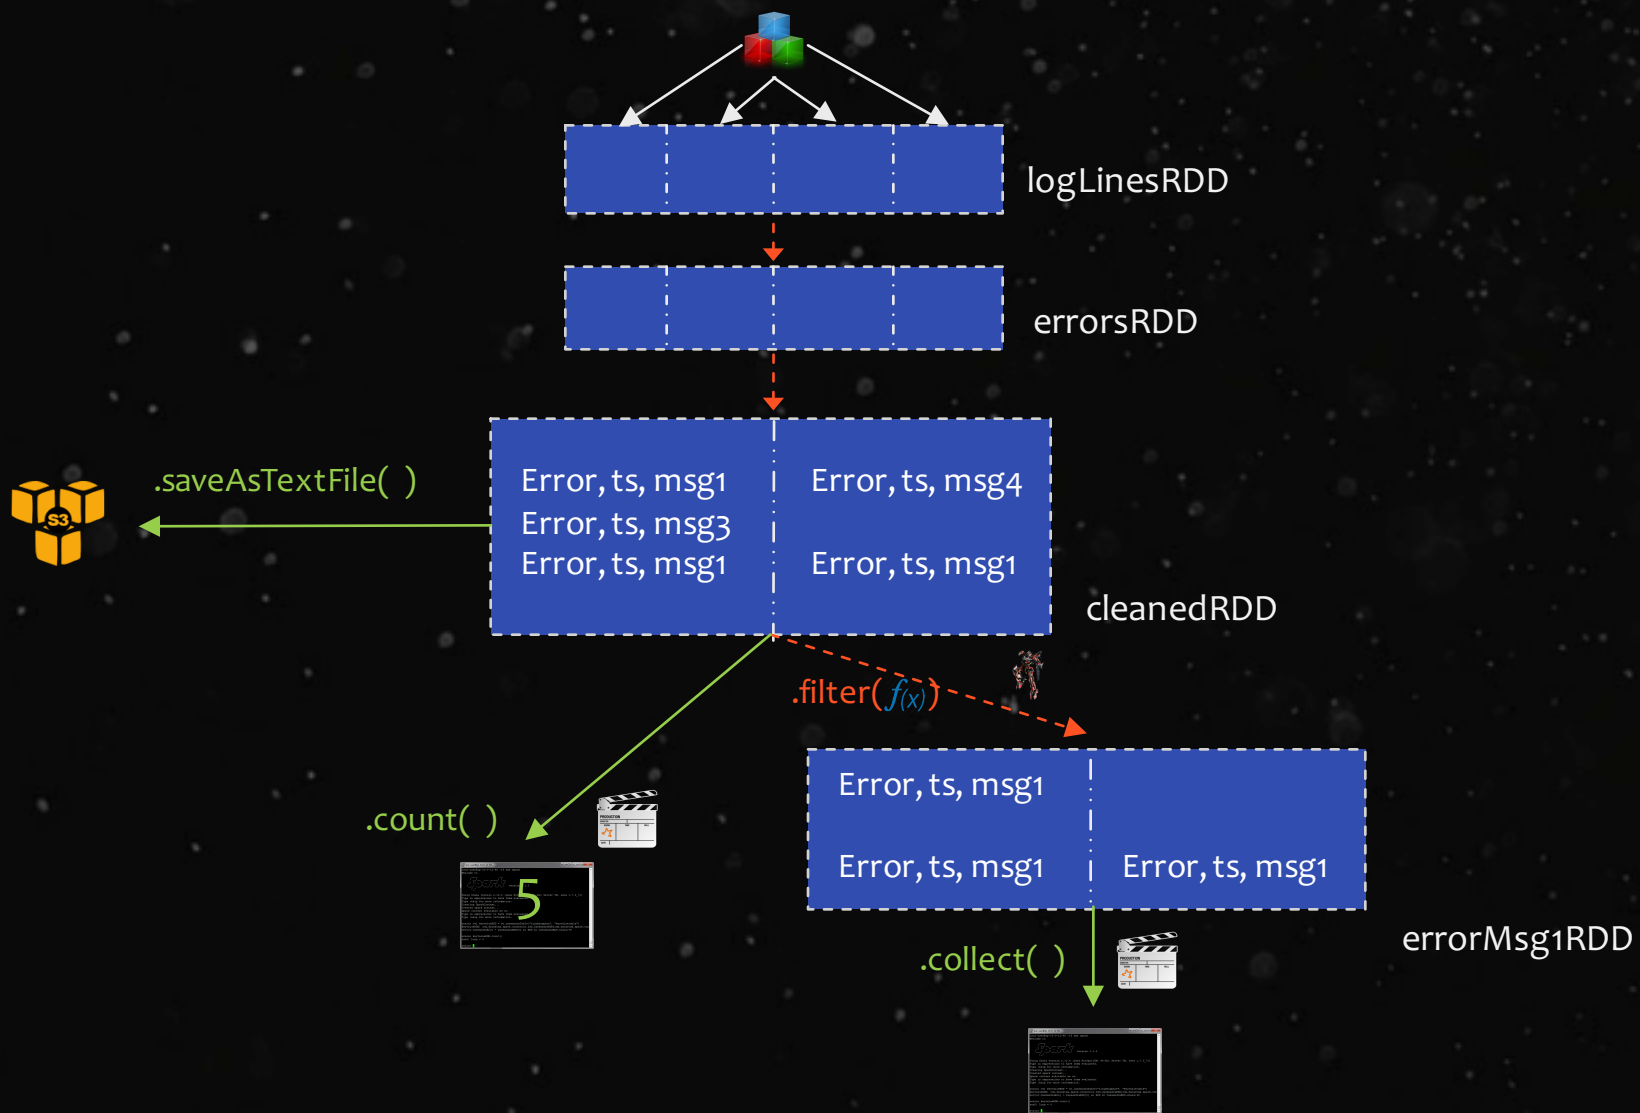
Driver

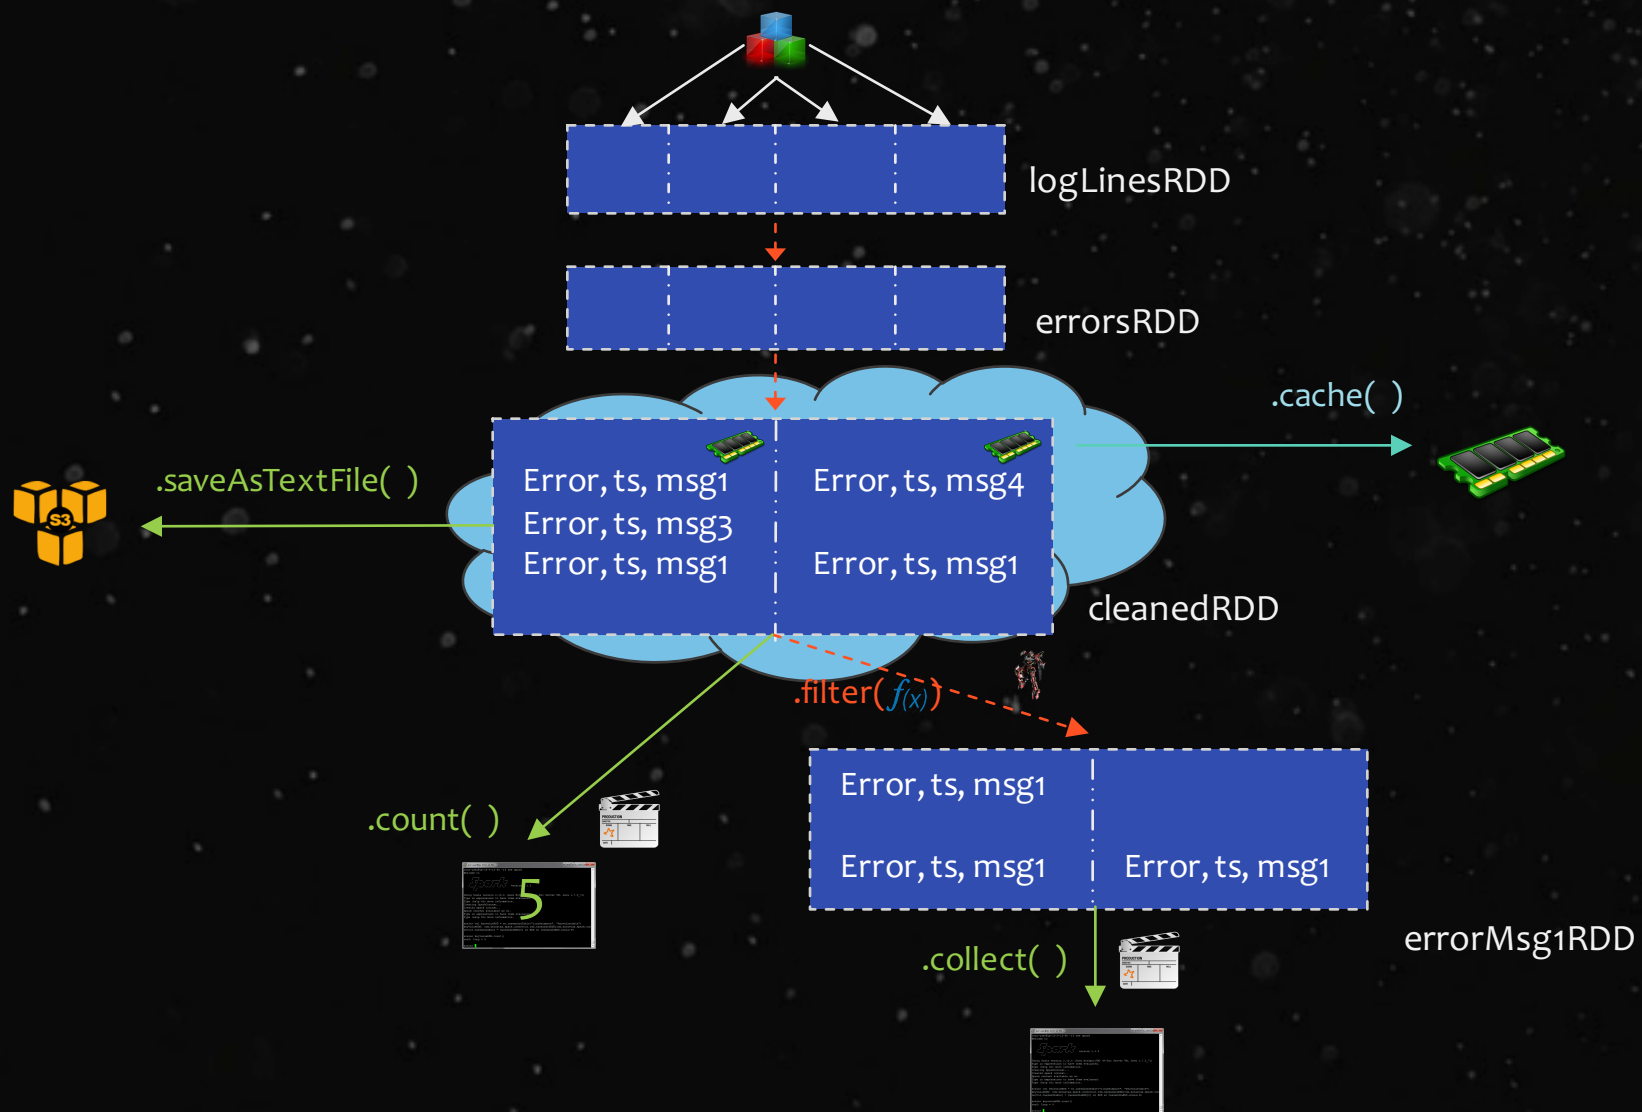















# PARTITION -> TASK -> PARTITION



# LIFECYCLE OF A SPARK PROGRAM

- 1) Create some input RDDs from external data or parallelize a collection in your driver program.
- 2) Lazily transform them to define new RDDs using transformations like `filter()` or `map()`
- 3) Ask Spark to `cache()` any intermediate RDDs that will need to be reused.
- 4) Launch actions such as `count()` and `collect()` to kick off a parallel computation, which is then optimized and executed by Spark.

# TRANSFORMATIONS (lazy)

<code>map()</code>	<code>intersection()</code>	<code>cartesion()</code>
<code>flatMap()</code>	<code>distinct()</code>	<code>pipe()</code>
<code>filter()</code>	<code>groupByKey()</code>	<code>coalesce()</code>
<code>mapPartitions()</code>	<code>reduceByKey()</code>	<code>repartition()</code>
<code>mapPartitionsWithIndex()</code>	<code>sortByKey()</code>	<code>partitionBy()</code>
<code>sample()</code>	<code>join()</code>	<code>...</code>
<code>union()</code>	<code>cogroup()</code>	<code>...</code>

# ACTIONS

`reduce()`

`collect()`

`count()`

`first()`

`take()`

`takeSample()`

`saveToCassandra()`

`takeOrdered()`

`saveAsTextFile()`

`saveAsSequenceFile()`

`saveAsObjectFile()`

`countByKey()`

`foreach()`

`...`

# TYPES OF RDDS

- HadoopRDD
- FilteredRDD
- MappedRDD
- PairRDD
- ShuffledRDD
- UnionRDD
- PythonRDD
- DoubleRDD
- JdbcRDD
- JsonRDD
- VertexRDD
- EdgeRDD
- **CassandraRDD** (*DataStax*)
- **GeoRDD** (*ESRI*)
- **EsSpark** (*ElasticSearch*)